ANUSHKA BABBAR

Milpitas, California | +1 650 200 8430 | anushka.sjsu@gmail.com | /anushkababbar | /anushka16121995

SUMMARY

Software Engineer with 3 years at Infosys building PySpark data pipelines and Kafka-driven backend systems. Now building production-style systems spanning event-driven backends, cloud data pipelines, and multi-agent AI applications, completing an MS in Engineering Management at San Jose State University.

EDUCATION

Master of Science in Engineering Management — Spring 2025 to Dec 2026 (Expected)
San Jose State University, California — GPA: 3.8/4.0

Post Graduate Diploma in Management — 2019 to 2021
Lal Bahadur Shastri Institute of Management, Delhi, India — GPA: 7.38/10

Bachelor of Technology in Information Technology — 2014 to 2018
Guru Gobind Singh Indraprastha University, Delhi, India — Percentage: 83.46%

TECHNICAL SKILLS

Languages: Java, Python, SQL **Testing:** JUnit 5, Mockito, MockMvc

Backend: Spring Boot, Spring JPA, Hibernate, Spring Security, REST APIs, MySQL, Microservices, System Design

Data Engineering: PySpark, Spark, Hive, Hadoop (HDFS), Parquet, Airflow, Pandas, NumPy, ETL Pipeline Design, Workflow Orchestration, Data Quality & Validation, Data Modeling & Schema Design, Kafka Streaming & Batch Ingestion

Cloud & Tools: GCP, Google Cloud Storage, BigQuery, Git, GitHub Actions, CI/CD, Docker, Postman

AI & LLM Tooling: Generative AI, Embeddings, RAG, Vector Databases (pgvector), Claude Code, Cursor, Cline

WORK EXPERIENCE

Infosys | Senior Associate Consultant — Bangalore, India | June 2021 to Feb 2024

- Cut data-pull requests by ~**30%** for **3+ downstream teams** by building PySpark ETL pipelines on Spark/Hadoop with batch and Kafka streaming ingestion into Hive tables.
- Achieved **zero data exposure incidents** on self-service financial analytics for Finance/Analytics teams by applying field-level encryption to PII in Hive and structuring schemas for safe joins on non-sensitive dimensions.
- Delivered **99.9%+ data accuracy** with **zero critical post-launch defects** by writing reconciliation scripts that validated record counts, balances, and schema consistency across SIT, UAT, and Go-Live.

PROJECTS

E-Commerce Order Processing System | Java, Spring Boot, Kafka, MySQL, PostgreSQL, pgvector, Docker

- Achieved **zero silent message loss** across **100+ REST endpoints** by building sync/async Kafka workflows with retry logic, a DLQ, and JWT-based role-based access control securing all role boundaries.
- Served similarity search on **768-dim** embeddings at ~**82ms** median latency by setting up PostgreSQL/pgvector, backed by **78% test coverage** and a CI pipeline.

Metadata-Driven Data Ingestion Platform | Python, PySpark, PostgreSQL, Hive, GCP/BigQuery, Docker, Airflow

- Processed **100K+** records at **99%+** pipeline success by building a Spring Boot metadata control plane tracking job/stage history in PostgreSQL, with a PySpark validator isolating **100%** of invalid records into quarantine.
- Detected stuck and failed jobs in real time by building job-status monitoring with automated failure notifications, extending the pipeline to GCP/BigQuery with Docker/Airflow orchestration.

Multi-Agent Assignment Grading System | Python, LLM Agents, Pandas, OpenPyXL, OCR/Vision, Prompt Engineering

- Cut manual grading time by **65%** on **40+** submissions per assignment by building an orchestrator coordinating Reader (Pandas/OpenPyXL) and OCR/Vision agents to extract and normalize mixed Excel/handwritten-PDF submissions.
- Matched **92%** instructor-grade scoring accuracy by building a Grader agent (rubric rules + LLM evaluation), with a Reviewer agent flagging **10%** of low-confidence grades for human review.